

MODEL EFFICIENCY & TRAINING

PARAMETER-EFFICIENT FINE-TUNING (PEFT) OF LARGE LANGUAGE MODELS

Adapting Large Models with Minimal Compute, Memory and Storage

Name: C. Anu Keerthana

Reg. No.: 192472057

Batch: BE-CSE-AI

1 • INTRODUCTION & PROBLEM

INTRODUCTION / BACKGROUND

Large Language Models (LLMs) contain billions of parameters. Adapting a pretrained LLM to a new task or domain traditionally means updating all of these parameters through full fine-tuning. Parameter-Efficient Fine-Tuning (PEFT) instead freezes most of the base model and trains only a small number of additional or selected parameters, making customization far cheaper and faster.

PROBLEM STATEMENT

Full fine-tuning of an LLM requires huge GPU memory, long training time and a full copy of the model for every task, which is impractical for most teams. Multiple tasks or clients need multiple fine-tuned models. The challenge is to specialize an LLM for new tasks or domains without retraining or storing the entire model each time.

2 • PEFT TECHNOLOGY

WHAT POWERS PEFT?

- Frozen pretrained transformer backbone
- Small set of trainable parameters injected
- Low-rank matrix decomposition (LoRA)
- Reparameterization / adapter modules
- Quantized base weights (QLoRA)

KEY PEFT METHODS

- LoRA / QLoRA – low-rank update matrices
- Adapter layers – small bottleneck modules
- Prefix / Prompt tuning – learnable tokens
- BitFit – bias-terms-only tuning
- IA3 – learned rescaling vectors

HOW IT WORKS

Base model weights stay frozen. A small number of new or selected parameters (often <1% of the model) are trained on task data. At inference, the small update is merged with or added alongside the frozen base model to produce task-specific behaviour.

3 • ARCHITECTURE / WORKFLOW / WORKING PRINCIPLE

01

BASE MODEL

Pretrained LLM, frozen

02

TASK DATA

Domain / task examples

03

INSERT ADAPTER

LoRA / adapter layers added

04

TRAIN SMALL SET

Only new params updated

05

MERGE / DEPLOY

Combine with base model

06

REUSE

Swap adapters per task

EFFICIENT TRAINING LOOP

Freeze base weights → attach lightweight trainable module → fine-tune on small task dataset → evaluate → merge or store adapter → reload a different adapter for the next task, reusing the same frozen base model throughout.

4 • INDUSTRY APPLICATIONS & REAL-WORLD USE CASE

Domain Adaptation

Specializing a general LLM for legal, medical or financial text

Multi-Task Deployment

One frozen base model + many swappable task adapters

On-Device / Edge AI

Small adapter files enable lightweight, low-memory deployment

Chatbot Customization

Brand voice, persona and support-policy tuning

Low-Resource Languages

Efficient adaptation where labeled data is limited

Rapid Prototyping

Fast, low-cost experimentation across many use cases

REAL-WORLD CASE | HUGGING FACE PEFT & QLoRA

Hugging Face's open-source PEFT library packages LoRA, prefix-tuning and related methods so teams can fine-tune large open models on a single GPU. QLoRA extends this by fine-tuning 4-bit quantized models, showing how PEFT lets organizations customize billion-parameter LLMs with modest hardware.

5 • IMPACT, RESPONSIBLE AI & FUTURE

ADVANTAGES

- Far lower compute and memory needs
- Small, easy-to-store adapter files
- Faster training and iteration
- One base model serves many tasks
- Reduced risk of catastrophic forgetting

CHALLENGES / LIMITATIONS

- Slight accuracy gap vs full fine-tuning
- Hyperparameter and rank sensitivity
- Adapter compatibility across model updates
- Tooling and versioning complexity
- Still inherits base model's biases/errors

FUTURE SCOPE

- Automated selection of PEFT method/rank
- Composable, stackable multi-adapters
- PEFT for multimodal foundation models
- Hardware-aware, on-device PEFT
- Standardized adapter-sharing hubs

CONCLUSION

PEFT reframes fine-tuning from "retrain everything" to "train only what's necessary." By freezing the base LLM and learning small, swappable modules such as LoRA adapters, PEFT makes large-model customization affordable, faster and more accessible — without sacrificing the knowledge already captured by pretraining.

REFERENCES / RESEARCH BASIS

1. Hu et al. (2021), "LoRA: Low-Rank Adaptation of Large Language Models."
2. Housby et al. (2019), "Parameter-Efficient Transfer Learning for NLP."
3. Detrmers et al. (2023), "QLoRA: Efficient Finetuning of Quantized LLMs."
4. Hugging Face, "PEFT Library Documentation" (2024-2026).